

# A Hybrid Approach for Reinforcement Learning Using Virtual Policy Gradient for Balancing an Inverted Pendulum

Doctoral Consortium

# Dylan Bates

Department of Mathematics  
Center for Research in Scientific Computation  
North Carolina State University

# Table of Contents

1. Research Problem

2. Outline of Objectives

3. State of the Art

4. Methodology

5. Expected Outcome

6. Stage of the Research



# Introduction

Reinforcement Learning is the process of using trial-and-error with finite rewards.

- Input state of the environment
- Output optimal action
- (Optional) Rewards

Our agent is an underactuated single-inverted pendulum on a one-dimensional track. The nonlinear equations of motion result in an unstable equilibrium that can be difficult to maintain.

# Environment and Agent



# Goals

- Use RL to train a realistic simulation of a virtual pole to balance itself.
- Use the trained neural network to balance a real pole in the real world.


Ultimately, training through simulation can speed up training time and increase robustness.

# Literature Review

- Traditional control: Linearized systems limit the angle of deflection. Robust performance is only achieved with properly tuned weights.
- Virtual RL: Usually implemented in a simulation like OpenAI's Cartpole. Physics here are usually questionable.
- Real RL: “Nothing works; I mean, the robots will break down - they'll break down all the time.” - Dr. Tim Lillicrap

# Environment

- Started with OpenAI Cartpole
- Adjusted physics, improved EOM
- Continuous action-space
- State:  $[x, \alpha, \dot{x}, \dot{\alpha}]$




**Figure:** Modified Cartpole environment.

# Neural Network

- Input state vector
- ReLU activation
- 2 outputs:  $\sim N(\mu, \sigma)$

Figure: Artificial NN approximates actions.





# Discounted Rewards

Rewards:  $r_t = 1, 0 \leq t \leq T$

$$R_t = \sum_{k=0}^{\infty} \gamma^k r_{t+k}$$

- $0 < \gamma < 1$  is a discount factor, intended to prioritize actions now over actions in the future.
- Larger values of  $\gamma$  takes a long term approach.

# Loss

$$L = -R \cdot \tilde{a} + \epsilon H$$

- $R$  = normalized discounted rewards
- $\tilde{a}$  = log probability of actions
- $H$  = optional entropy parameter

This can be modified to promote optimal behaviour.

# Neural Networks - Policy Gradient

Unlike supervised learning, we cannot calculate an explicit error between the neural network's output and the “correct” answer. Goal:  $\max \sum_i \ln p(y_i|x_i)$

- Maximize expected return:

$$J(\pi_\theta) = \int_{\tau} P(\tau|\theta) R(\tau) = \mathbb{E}_{\tau \sim \pi_\theta} [R(\tau)]$$

- Use gradient descent:

$$\nabla_{\theta} J(\pi_\theta) = \mathbb{E}_{\tau \sim \pi_\theta} [\nabla_{\theta} \ln P(\tau|\theta) R(\tau)]$$

- Normalized discounted rewards encourage and discourage half of actions.

# Virtual Policy Gradient

$$\nabla_{\theta} J(\pi_{\theta}) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[ \sum_{t=0}^T \nabla_{\theta} \ln \pi_{\theta}(\mathbf{a}_t | \mathbf{s}_t) A^{\pi_{\theta}}(\mathbf{s}_t, \mathbf{a}_t) \right]$$

The algorithm is incredibly generalizeable, allowing the same code to complete a variety of tasks.

- Balance a pole
- Play Pong
- Drive a car

You only need an updated simulation, and the appropriate number of neurons, hyperparameters...

# Results

The inverted pendulum learned to balance for at least 10 seconds 91% of the time

- takes an average of 807 trials
- high variance among gradients and training time
- optimizing hyperparameters reduced this to 355
- minimum trials was only 54 ( $\sim 1$  minute to train)

Due to the high variance among the gradients, even the best hyperparameters had inconsistent performance.

# New Work - Actor Critic and PPO

- Actor: ANN predicting the best action to take
- Critic: new head of the same ANN estimating the value of that action,  $V^\pi(s_t)$
- Subtracting this baseline reduces the variance of Policy Gradient

Given  $Q^\pi(s, a) = \mathbb{E}_{\tau \sim \pi} [R(\tau) | s_0 = s, a_0 = a]$ ,  
set  $A^\pi(a_t | s_t) = Q^\pi(s_t, a_t) - V^\pi(s_t)$

Then we use the same equation as before:

$$\nabla_{\theta} J(\pi_{\theta}) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[ \sum_{t=0}^T \nabla_{\theta} \ln \pi_{\theta}(a_t | s_t) A^{\pi_{\theta}}(s_t, a_t) \right]$$

# Applications

Ultimately, training in simulation could make other machine learning applications more efficient, speeding up the development of control systems that can be implemented in the real world.

- Self-driving cars
- Reusable rockets
- Complex robotics
- Healthcare treatment

Domain randomization and cross-modal learning can help make virtually trained models more robust.

# Questions and Contact Info

- Email:  
`dwbates@ncsu.edu`
- LinkedIn:  
`linkedin//dwgb93`
- Videos/Slides/Paper:  
`https://go.ncsu.edu/icaart-videos`
- Happy Birthday, Mom!

